

# Caches and the Memory Hierarchy

# Latency, size and price of computer memory

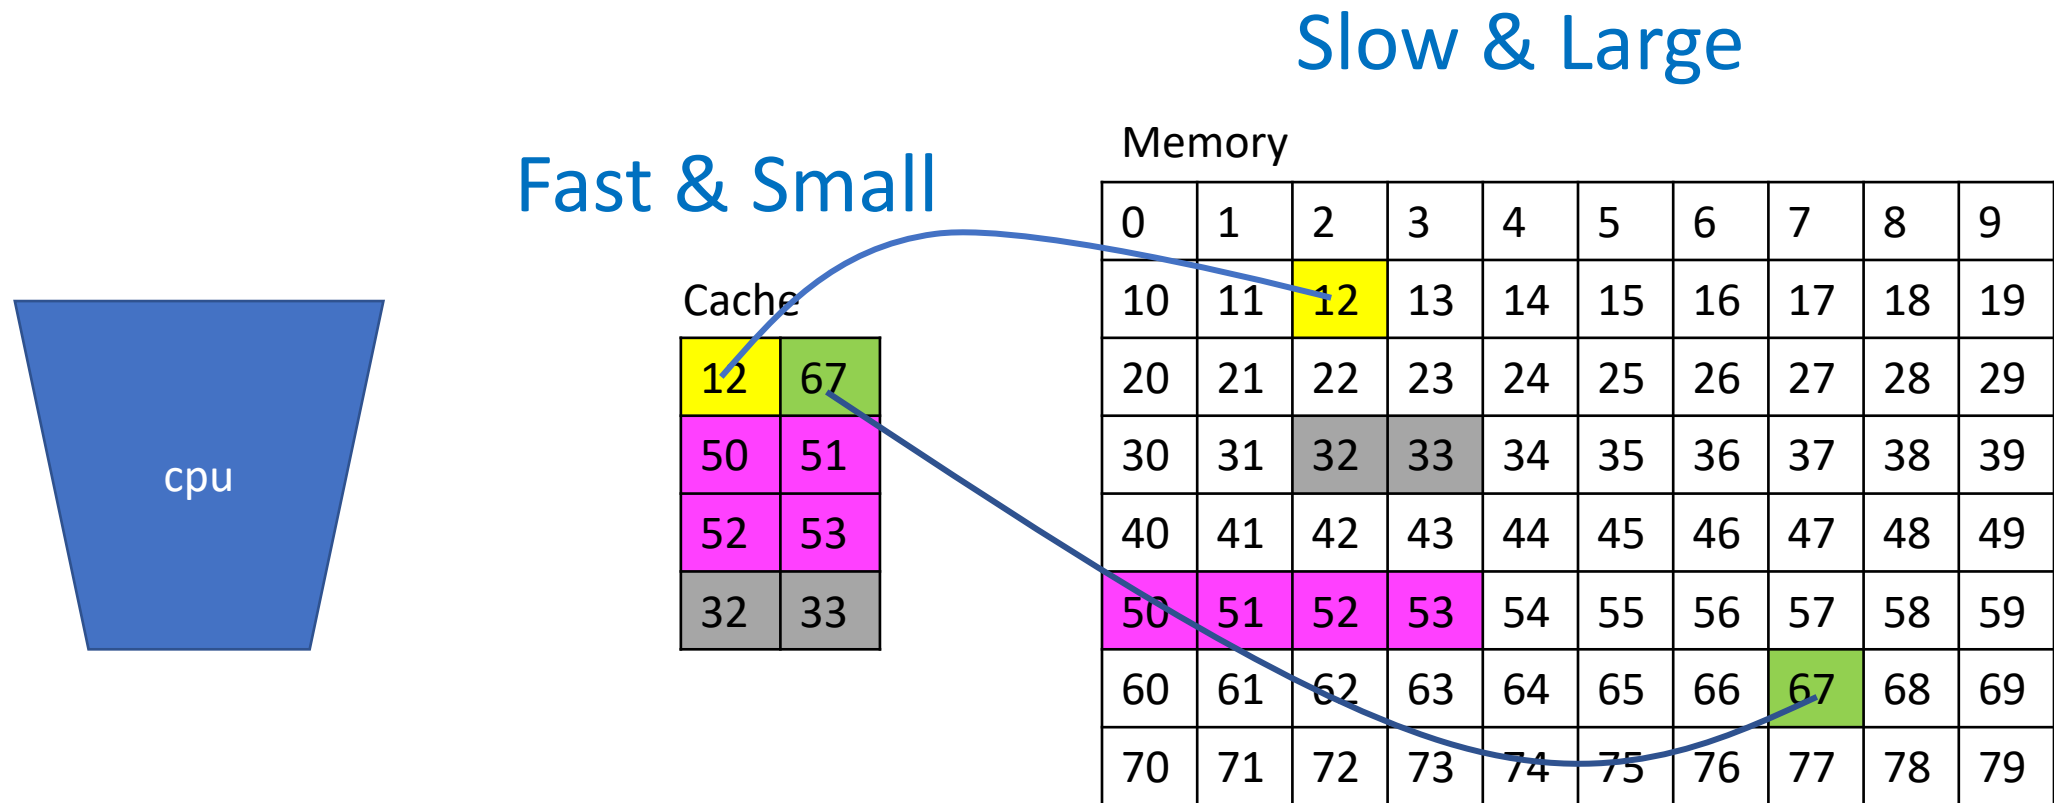
Given a budget, we need to trade off

## \$10: Fast & Small

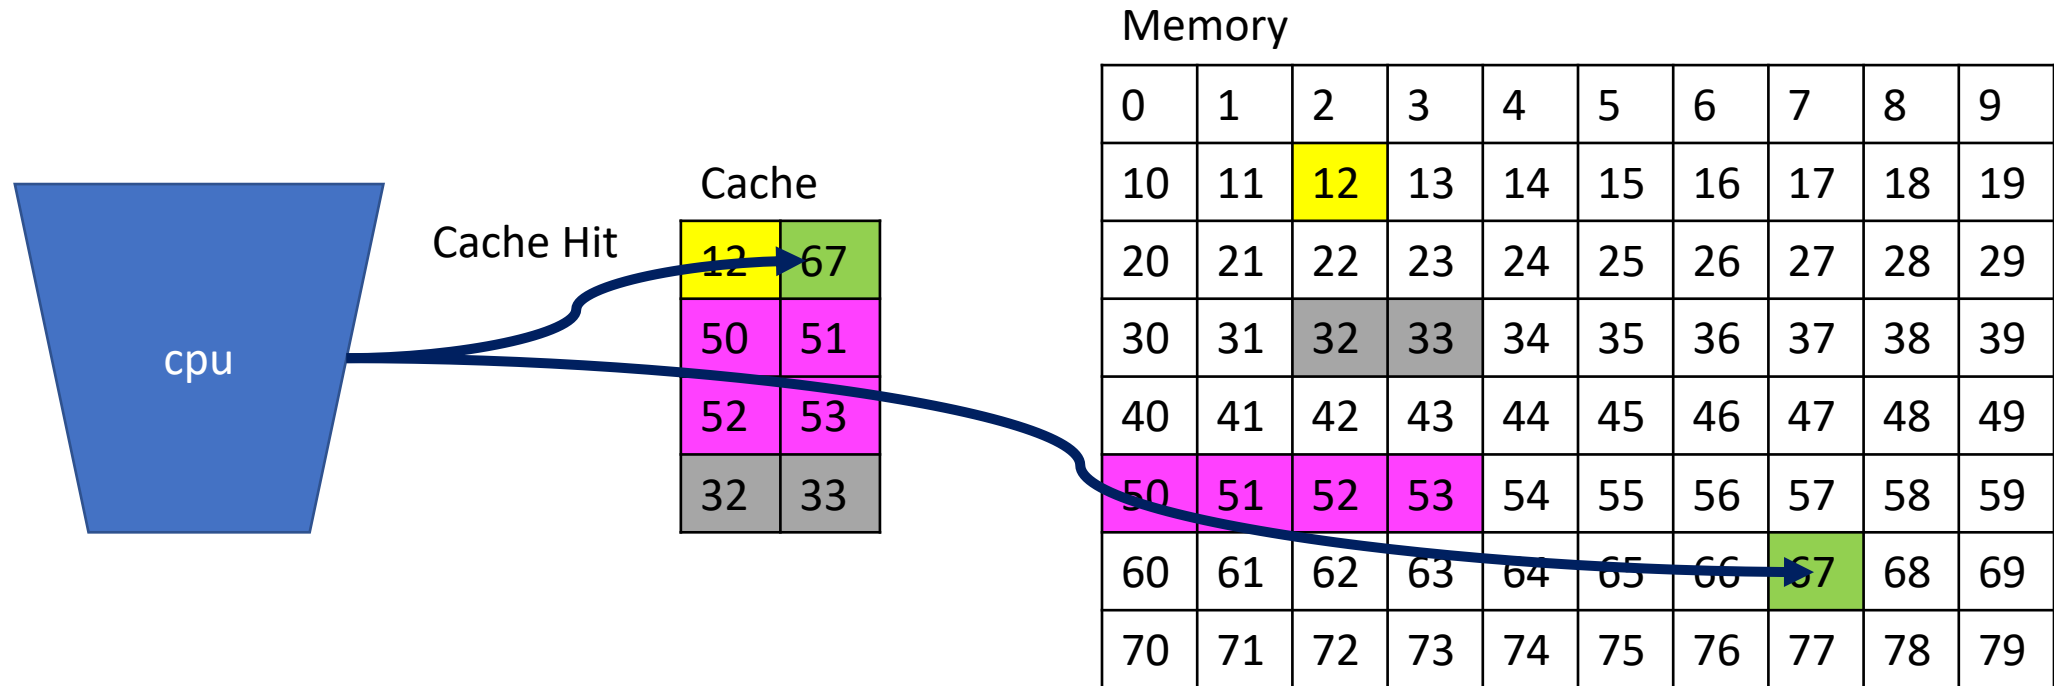

## \$10: Slow & Large

[illegible]

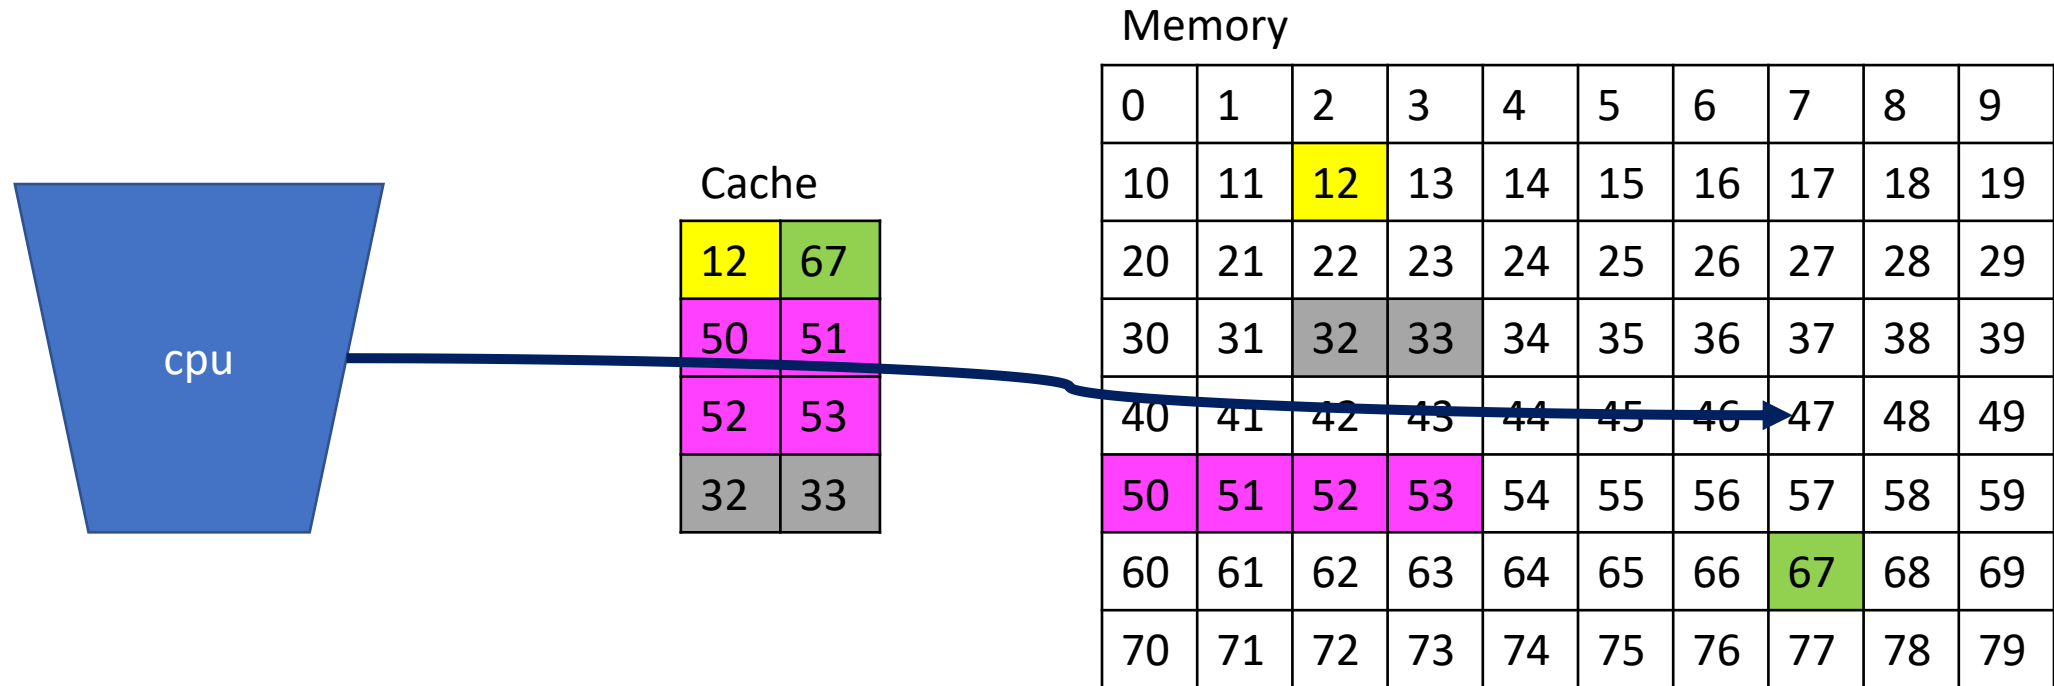
# Cache: The basic idea



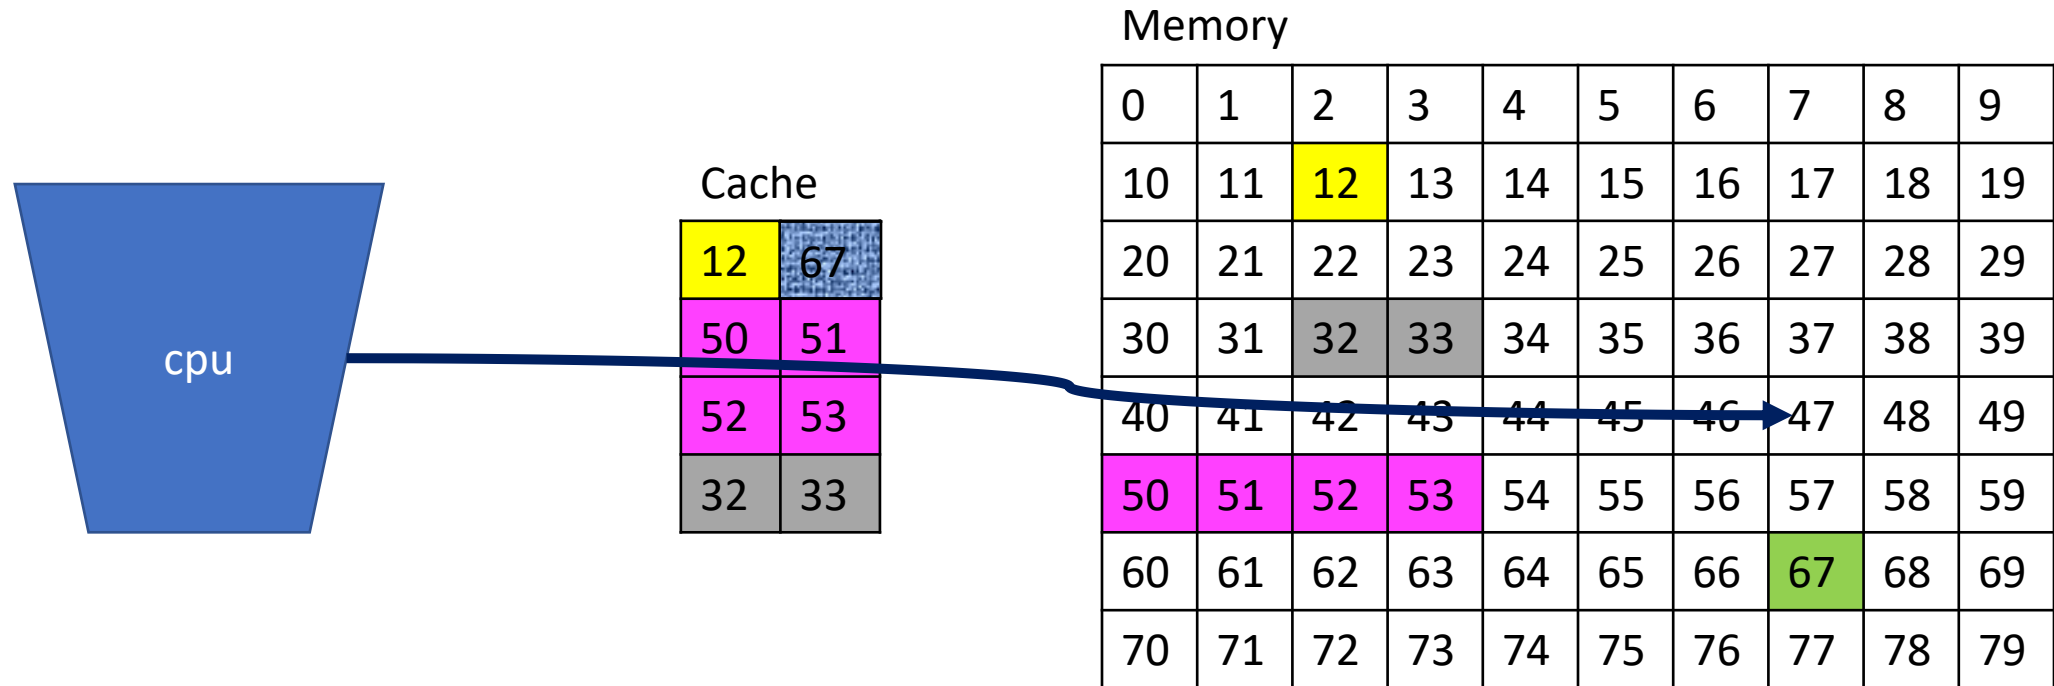
# Cache Hit



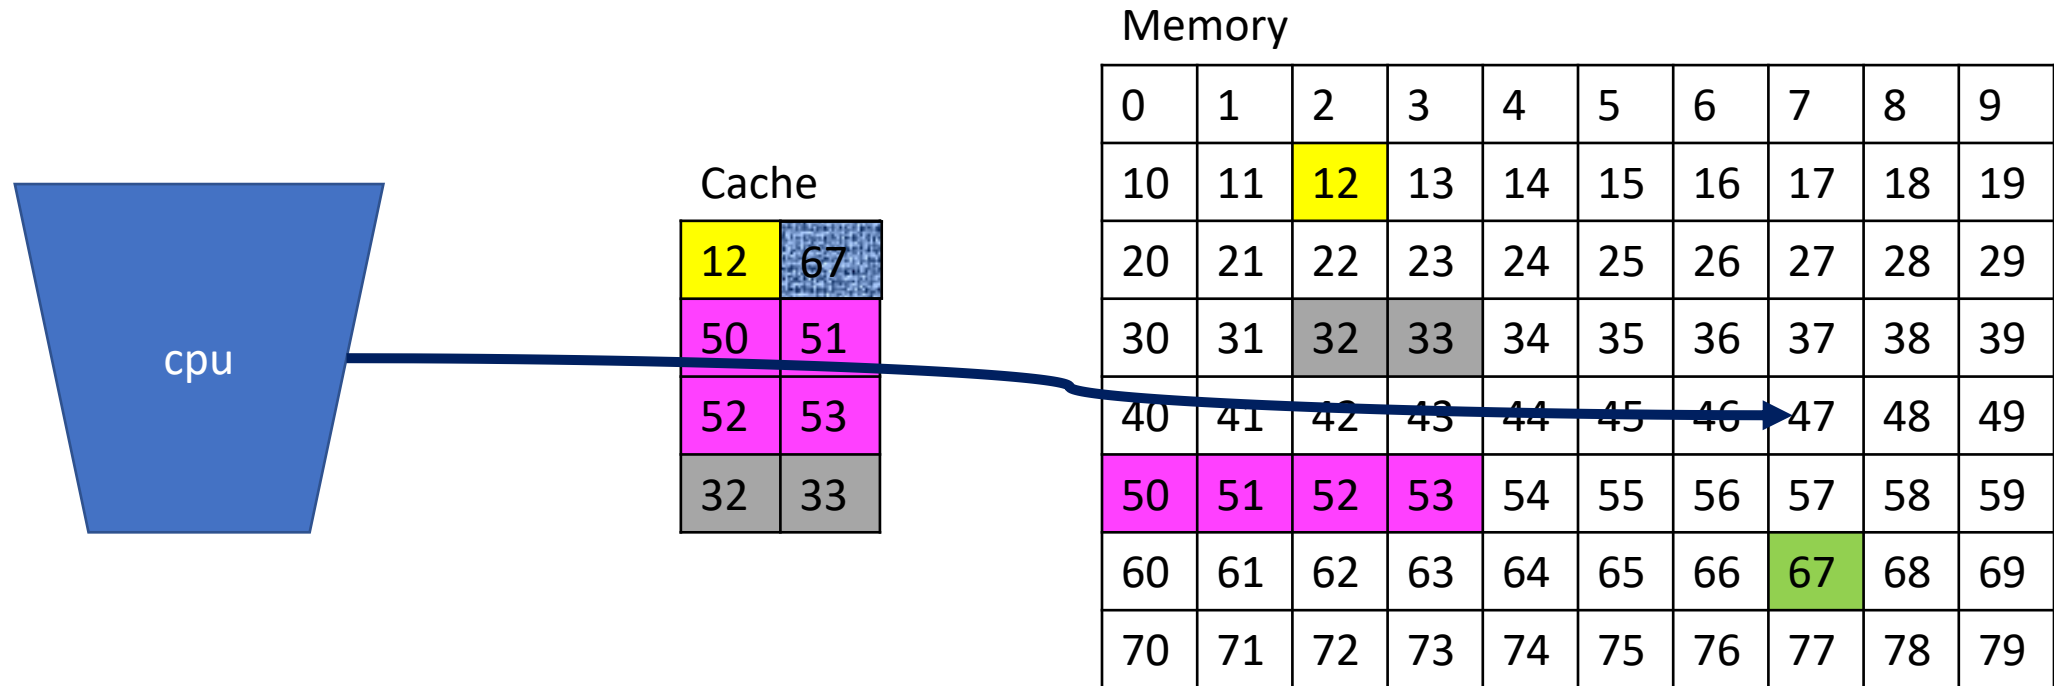
# Cache Miss



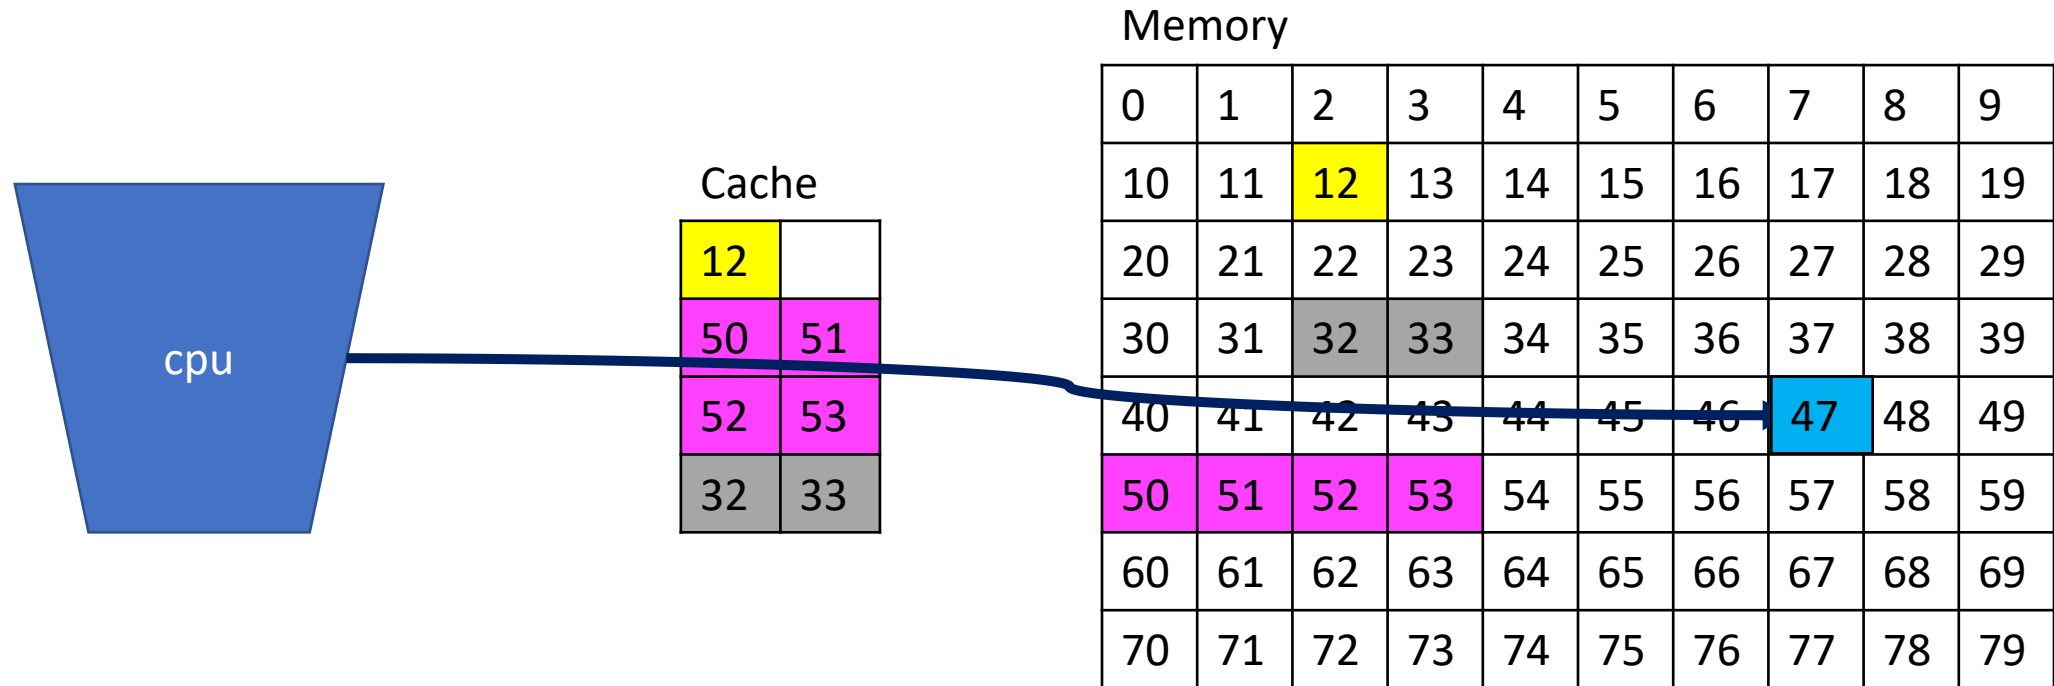
# Cache Miss Service: 1) Choose byte to drop



# Cache Miss Service: 2) write back



# Cache Miss Service: 3) Read In





# Access Locality

- The cache is effective If most accesses are hits.
  - Cache Hit Rate is high.
- **Temporal Locality**: Multiple accesses to **same** address within a short time period

# Spatial locality

- **Spatial Locality**: Multiple accesses to close-together addresses in short time period.
  - The difference between two sums.
  - Counting words by sorting
- Benefiting from spatial locality
  - Memory is partitioned into **Blocks/Lines** rather than single bytes.
  - Moving a block of memory takes much less time than moving each byte individually.
  - Memory locations that are close to each other are likely to fall in the same block.
  - Resulting in more cache hits.

# Unsorted word count / poor locality

=== unsorted list:

the, vernacular, but, as, for, you, ye, carrion, rogues, turning, to,

- Consider the memory access to the dictionary D:
- Count without sort:  
D[the]=12332,...,D[but]=943,.....,D[vernacular]=10,.....,D[for]=..
- Temporal locality for very common words like “the”
- No spatial locality

# sorted word count / good locality

```
=== sorted list:
```

```
lines, lingered, lingered, lingered, lingered, lingered, lingering,  
lingering, lingering, lingering, lingering, lingering, lingering,  
lingering, lingers, lingo, lingo, lining, link, link, linked, linked,  
linked, linked, links, links
```

Entries to D are added one at a time.

1. D[lines]=33
2. D[lines]=33, D[lingered]=5
3. D[lines]=33, D[lingered]=5, D[lingering]=8

Assuming new entries are added at the end, this gives spatial locality.

Spatial locality makes code run much faster (X300)

# Summary

- Caching reduces storage latency by bringing relevant data close to the CPU.
- This requires that code exhibits access locality:
  - Temporal locality: Accessing the same location multiple times
  - Spatial locality: Accessing neighboring locations.